

TensorRT的INT8量化原理

- **Goal:** Convert FP32 CNNs into INT8 without significant accuracy loss.
- **Why:** INT8 math has higher throughput, and lower memory requirements.
- **Challenge:** INT8 has significantly lower precision and dynamic range than FP32.
- **Solution:** Minimize loss of information when quantizing trained model weights to INT8 and during INT8 computation of activations.
- **Result:** Method was implemented in TensorRT. It does not require any additional fine tuning or retraining.

INT8 Inference

Challenge

- INT8 has significantly **lower precision and dynamic range** compared to FP32.

| | Dynamic Range | Min Positive Value |
|------|--|-----------------------|
| FP32 | $-3.4 \times 10^{38} \sim +3.4 \times 10^{38}$ | 1.4×10^{-45} |
| FP16 | $-65504 \sim +65504$ | 5.96×10^{-8} |
| INT8 | $-128 \sim +127$ | 1 |

- Requires more than a simple type conversion from FP32 to INT8.

Linear quantization

Representation:

Tensor Values = FP32 scale factor * int8 array + FP32 bias



Symmetric linear quantization


Representation:

Tensor Values = FP32 scale factor * int8 array

One FP32 scale factor for the entire int8 tensor


Quantization

- No saturation: map $|\max|$ to 127



- Significant accuracy loss, in general




- Saturate above $|\text{threshold}|$ to 127



- Weights: no accuracy improvement
- Activations: improved accuracy
- Which $|\text{threshold}|$ is optimal?

Q: How to optimize threshold selection?

- It's always a tradeoff between **range** and **precision** of the INT8 representation.




A: Minimize information loss, since FP32 \rightarrow INT8 is just re-encoding information.

“Relative Entropy” of two encodings

- INT8 model encodes the **same information** as the original FP32 model.
- We want to **minimize loss of information**.
- Loss of information is measured by Kullback-Leibler divergence (AKA *relative entropy* or *information divergence*).
 - P, Q - two discrete probability distributions.
 - $\text{KL_divergence}(P, Q) := \text{SUM}(P[i] * \log(P[i] / Q[i]), i)$
- **Intuition**: KL divergence measures the amount of information lost when approximating a given encoding.

Solution: Calibration



- Run FP32 inference on Calibration Dataset.
- For each Layer:
 - collect histograms of activations.
 - generate many quantized distributions with different saturation thresholds.
 - pick threshold which minimizes $KL_divergence(ref_distr, quant_distr)$.
- Entire process takes a few minutes on a typical desktop workstation.

如何寻找最优的阈值T使得精度的损失最小？

NVIDIA选择的是KL-divergence，其实就是相对熵。相对熵表述的就是两个分布的差异程度，这里就是量化前后两个分布的差异程度。差异最小就是最好的了，因此问题转换为求相对熵的最小值。

$$D_{KL}(p||q) = \sum_{i=1}^N p(x_i) \cdot (\log p(x_i) - \log q(x_i))$$

KL散度来精确测量这种最优和次优之间的差异。

F32就是原来的最优编码，INT8就是次优的编码，用KL散度来描述这两种编码之间的差异。

相对熵表示的是采用次优编码时会多需要多少个bits来编码，也就是与最优编码之间的bit差；

而交叉熵表示的是用次优编码方式时确切需要多少个bits来表示；

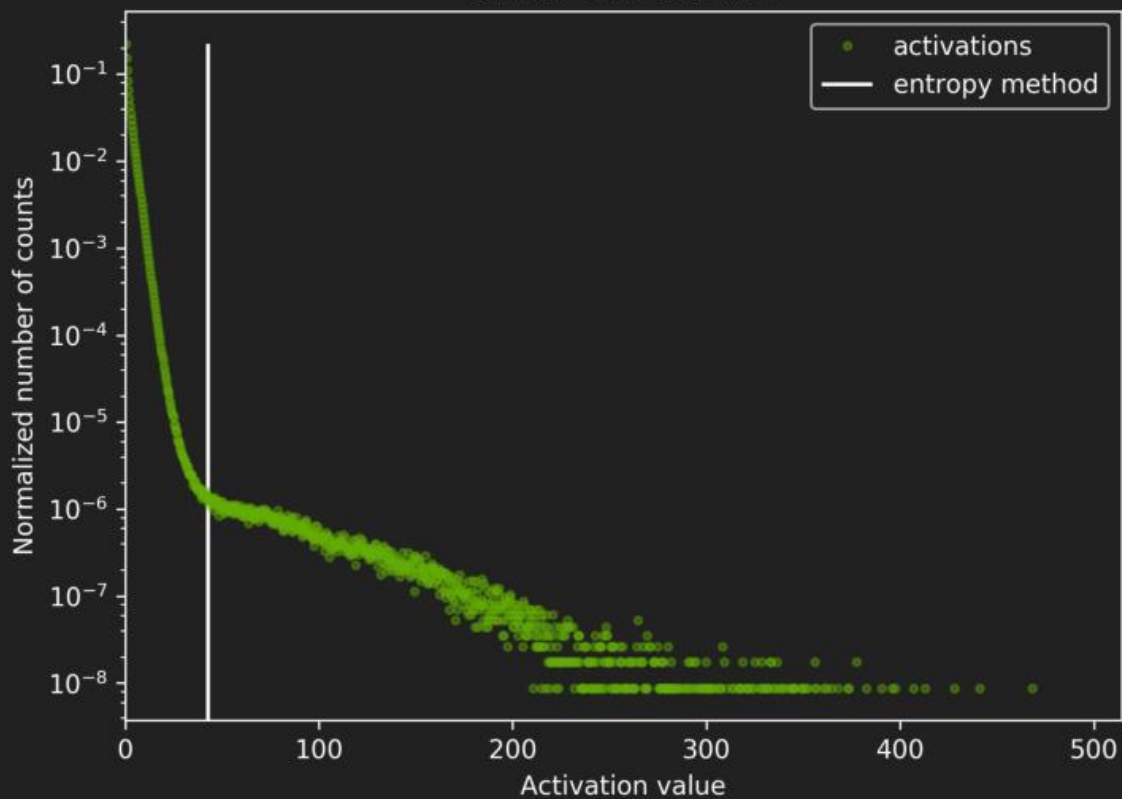
因此，最优编码所需要的bits=交叉熵-相对熵。

Typical workflow in TensorRT

- You will need:
 - Model trained in FP32.
 - Calibration dataset.
- TensorRT will:
 - Run inference in FP32 on calibration dataset.
 - Collect required statistics.
 - Run calibration algorithm → optimal scaling factors.
 - Quantize FP32 weights → INT8.
 - Generate “CalibrationTable” and INT8 execution engine.


Before saturation

resnet-152: res4b30



After saturation

resnet-152: res4b30



校准算法：

calibration：基于实验的迭代搜索阈值。

- 提供一个样本数据集（最好是验证集的子集），称为“校准数据集”，用来做校准。
- 在校准数据集上运行FP32推理。收集激活的直方图，并生成一组具有不同阈值的8位表示法，并选择具有最少KL散度的表示。

KL散度是在参考分布（即FP32激活）和量化分布之间（即8位量化激活）之间。

TRT提供了IInt8EntropyCalibrator，该接口需要由客户端实现，以提供校准数据集和一些用于缓存校准结果的样板代码。

Entropy Calibration - pseudocode

Input: FP32 histogram H with 2048 bins: bin[0], ..., bin[2047]

For i in range(128 , 2048):

```
reference_distribution_P = [ bin[ 0 ], ..., bin[ i-1 ] ]           // take first ' i ' bins from H
outliers_count = sum( bin[ i ], bin[ i+1 ], ... , bin[ 2047 ] )
reference_distribution_P[ i-1 ] += outliers_count
P /= sum(P)                                                       // normalize distribution P
candidate_distribution_Q = quantize [ bin[ 0 ], ..., bin[ i-1 ] ] into 128 levels // explained later
expand candidate_distribution_Q to ' i ' bins                      // explained later
Q /= sum(Q)                                                       // normalize distribution Q
divergence[ i ] = KL_divergence( reference_distribution_P, candidate_distribution_Q)
```

End For

Find index 'm' for which divergence[m] is minimal

threshold = (m + 0.5) * (width of a bin)

上面就是一个循环，不断地构造P和Q，并计算相对熵，然后找到最小（截断长度为m）的相对熵，此时表示Q能比较好地拟合P分布了。而阈值就等于 $(m + 0.5) \times \text{width of a bin}$ 。

要做INT8量化，需要：

- 原来的未量化的模型
- 一个校准数据集
- 进行量化过程的校准器

校准过程我们是不用参与的，全部都由TensorRT内部完成，但是，需要告诉校准器如何获取一个batch的数据，也就是说，需要重写校准器类中的一些方法。

- 准备一个校准集，用于在转换过程中寻找使得转换后的激活值分布与原来的FP32类型的激活值分布差异最小的阈值；
- 写一个校准器类，该类需继承`trt.IInt8EntropyCalibrator2`父类，并重写`get_batch_size`, `get_batch`, `read_calibration_cache`, `write_calibration_cache`这几个方法。
- 使用时，需额外指定`cache_file`，该参数是校准集cache文件的路径，会在校准过程中生成，方便下一次校准时快速提取。

Check if Your GPU Supports FP16/INT8

1. check your GPU Compute Capability

visit <https://developer.nvidia.com/cuda-gpus#compute> and check your GPU compute capability.

For example, GTX1080 is 6.1, Tesla T4 is 7.5.

2. check the hardware-precision-matrix

visit <https://docs.nvidia.com/deeplearning/tensorrt/support-matrix/index.html#hardware-precision-matrix> and check the matrix.

For example, compute capability 6.1 supports FP32 and INT8. 7.5 supports FP32, FP16, INT8, FP16 tensor core, etc.